

NATURAL LANGUAGE PROCESSING

MACHINE LEARNING - BASIC ALGORITHMS

DR. SAMPATH DEEGALLA

PHD, PHLIC (STOCKHOLM UNIVERSITY, SWEDEN),

BSc ENG (UNIVERSITY OF PERADENIYA), AMIESL, MIEEE, MIEEE-CS

sampathd@sltc.ac.lk

INPUT

OUTLINE

- Basic types of Learning in Data Mining
- Types of input to the learning algorithm
 - ▶ Concept, Example, Attribute
- Preparing the input
 - ▶ Missing values, Inaccurate values, getting to know data

BASIC TYPES OF LEARNING

- **Classification** (Supervised Learning):
 - ▶ predicting a particular class value (discrete)
 - ▶ E.g. Weather problem
- **Association learning**:
 - ▶ detecting associations between features
- **Clustering** (Unsupervised Learning):
 - ▶ grouping similar instances into clusters
- **Numeric prediction**:
 - ▶ predicting a numeric quantity
 - ▶ E.g. CPU performance

CLASSIFICATION LEARNING

- Example
 - ▶ Weather data to predict **play/not play**
- Classification learning is **supervised**
 - ▶ Scheme is being provided with actual outcome
 - ▶ Outcome is called the **class** of the example
- success can be measured on fresh data for which class labels are known (test data)

3

112

ASSOCIATION LEARNING

- Example
 - ▶ supermarket basket analysis - what items are bought together (e.g. milk + cereal)
- Can be applied if no class is specified and any kind of structure is considered “interesting”
- Difference with classification learning:
 - ▶ Can predict any attribute’s value, not just the class, and more than one attribute’s value at a time
 - ▶ Hence: far more association rules than classification rules
 - ▶ Thus: constraints are necessary
 - Minimum coverage and minimum accuracy

4

112

CLUSTERING

- Example
 - ▶ customer grouping
- Finding groups of items that are similar
- Clustering is **unsupervised**
 - ▶ The class of an example is not known
- Success often measured subjectively

5

112

NUMERIC PREDICTION

- Classification learning, but “class” is numeric
- Learning is supervised
 - ▶ Scheme is being provided with target value
- Measure success on test data

Outlook	Temperature	Humidity	Windy	Windy
Sunny	Hot	High	False	5
Sunny	Hot	High	True	0
Overcast	Hot	High	False	55
Rainy	Mild	Normal	False	40
...

6

112

WHAT'S A CONCEPT?

- Concept: thing to be learned
- Concept description: output of learning scheme
 - ▶ In the form of rule set, decision tree, ...
- Success rate on test data → How well the concept has been learned

7

112

WHAT'S IN AN EXAMPLE?

- Instance: specific type of example
 - ▶ Thing to be classified, associated, or clustered
 - ▶ Individual, independent example of target concept
 - ▶ Characterized by a predetermined set of attributes
- Input to learning scheme: set of instances/dataset
 - ▶ Represented as a single relation/flat file

8

112

GENERATING A FLAT FILE

- Process of flattening a file is called **denormalization**
 - ▶ Several relations are joined together to make one
- Possible with any finite set of finite relations
- Denormalization may produce fake regularities that reflect structure of database
 - ▶ Example: "supplier" predicts "supplier address"

9

112

WHAT'S IN AN ATTRIBUTE?

- Each instance is described by a fixed predefined set of features, its "attributes"
- But: number of attributes may vary in practice
 - ▶ Possible solution: "irrelevant value" flag
- Related problem: existence of an attribute may depend of value of another one
 - ▶ E.g. spouse name → married/single
- Possible attribute types :
 - ▶ Nominal, ordinal, interval and ratio

10

112

NOMINAL QUANTITIES

- Values are distinct symbols
 - ▶ Values themselves serve only as labels or names
 - ▶ Nominal comes from the Latin word for name
- Example: attribute “outlook” from weather data
 - ▶ Values: “sunny”, “overcast”, and “rainy”
- No relation is implied among nominal values
 - ▶ no ordering or distance measure
- Only equality tests can be performed

11

112

ORDINAL QUANTITIES

- Impose order on values
- But: no distance between values defined
- Example: attribute “temperature” in weather data
 - ▶ Values: “hot” > “mild” > “cool”
- Note: addition and subtraction don’t make sense
- Example rule:
 - ▶ $\text{temperature} < \text{hot} \rightarrow \text{play} = \text{yes}$
- Distinction between nominal and ordinal not always clear (e.g. attribute “outlook”)
- Note: Distinction between nominal and ordinal is not always straightforward

12

112

INTERVAL QUANTITIES (NUMERIC)

- Interval quantities are not only ordered but measured in fixed and equal units
- Example 1: attribute “temperature” expressed in degrees Fahrenheit
- Example 2: attribute “year”
- Difference of two values makes sense
 - ▶ E.g. 46F, 48F and 1939, 1945
- Sum or product doesn’t make sense
 - ▶ Zero point is not defined! (year 0)

13

112

RATIO QUANTITIES

- Ratio quantities are ones for which the measurement scheme defines a zero point
- Example: attribute “distance”
 - ▶ Distance between an object and itself is zero
- Ratio quantities are treated as real numbers
 - ▶ All mathematical operations are allowed
- But: is there an “inherently” defined zero point?
 - ▶ Answer depends on scientific knowledge (e.g. Fahrenheit knew no lower limit to temperature)

14

112

ATTRIBUTE TYPES USED IN PRACTICE

- Most schemes accommodate just two levels of measurement: nominal and ordinal
- Nominal attributes are also called “categorical”, “enumerated”, or “discrete”
- Special case: dichotomy (“boolean” attribute)
- Ordinal attributes are called “numeric”, or “continuous”
 - ▶ But: “continuous” implies mathematical continuity

15

112

PREPARING THE INPUT

- Problem: different data sources (e.g. sales department, customer billing department, ...)
 - ▶ Differences: styles of record keeping, conventions, time periods, data aggregation, primary keys, errors
 - ▶ Data must be assembled, integrated, cleaned up
 - ▶ “Data warehouse”: consistent point of access
- Denormalization relational data are not the only issue
- External data may be required (“overlay data”)
- Critical: type and level of data aggregation

16

112

THE WEATHER DATA

Outlook	Temperature	Humidity	Windy	Play
sunny	hot	high	false	no
sunny	hot	high	true	no
overcast	hot	high	false	yes
...				

- Instances, Examples, Objects – Represent the rows of the table.
- Attributes, Features – Represent the columns ‘Outlook’, ‘Temperature’, ‘Humidity’, and ‘Windy’.
- Class Attribute – Represented by the column ‘Play’.

17

112

MISSING VALUES

- Frequently indicated by out-of-range entries
 - Types: unknown, unrecorded, irrelevant
 - Reasons:
 - malfunctioning equipment
 - changes in experimental design
 - collation of different datasets
 - measurement not possible
- Missing value may have significance in itself (e.g., missing test in a medical examination)
 - Most schemes assume that is not the case
 - ⇒ “missing” may need to be coded as an additional value

18

112

INACCURATE VALUES

- Reason: data has not been collected for mining it
- Result: errors and omissions that don't affect original purpose of data (e.g., age of customer)
- Typographical errors in nominal attributes \Rightarrow values need to be checked for consistency
- Typographical and measurement errors in numeric attributes \Rightarrow outliers need to be identified
- Errors may be deliberate (e.g., wrong zip/postal codes)
- Other problems: duplicates, scale data

19

112

GETTING TO KNOW THE DATA

- Simple visualization tools are very useful
 - ▶ Nominal attributes: histograms (Distribution consistent with background knowledge?)
 - ▶ Numeric attributes: graphs (Any obvious outliers?)
- 2-D and 3-D plots show dependencies
- Need to consult domain experts
- Too much data to inspect? Take a sample!

20

112

BASIC ALGORITHMS

OUTLINE

- Basic Algorithms
 - ▶ IR
 - ▶ Decision Trees
 - ▶ Decision Rules (PRISM)
 - ▶ Nearest Neighbor
 - ▶ Naïve Bayes

21

112

ONER

SIMPLICITY FIRST

- Simple algorithms often work very well!
- There are many kinds of simple structure, eg:
 - ▶ One attribute does all the work
 - ▶ All attributes contribute equally & independently
 - ▶ A weighted linear combination might do
 - ▶ Instance-based: use a few prototypes
 - ▶ Use simple logical rules
- Success of method depends on the domain

22

112

INFERRING RUDIMENTARY RULES

- 1R: learns a 1-level decision tree
 - ▶ i.e., rules that all test one particular attribute
- Basic version
 - ▶ One branch for each value
 - ▶ Each branch assigns most frequent class
 - ▶ Error rate: proportion of instances that don't belong to the majority class of their corresponding branch
 - ▶ Choose attribute with lowest error rate (assumes nominal attributes)

23

112

ZERO R AND ONE R ALGORITHMS

- **ZeroR**: Outputs the majority class.
- **OneR**: Learns a 1-level decision tree.
 - ▶ Rules that test one particular attribute.
 - ▶ Basic version:
 - One branch for each attribute value.
 - Each branch assigns the most frequent class.
 - ▶ Error rate is the proportion of instances that don't belong to the majority class of their corresponding branch.
 - ▶ Choose the attribute with the lowest error rate.
- The basic version of OneR assumes nominal attributes.

24

112

THE WEATHER DATA

Outlook	Temperature	Humidity	Windy	Play
sunny	hot	high	false	no
sunny	hot	high	true	no
overcast	hot	high	false	yes
rainy	mild	high	false	yes
rainy	cool	normal	false	yes
rainy	cool	normal	true	no
overcast	cool	normal	true	yes
sunny	mild	high	false	no
sunny	cool	normal	false	yes
rainy	mild	normal	false	yes
sunny	mild	normal	true	yes
overcast	mild	high	true	yes
overcast	hot	normal	false	yes
rainy	mild	high	true	no

25

112

OUTPUT OF ZEROR

Outlook	Temp	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
...				

Play → Yes

26

112

PSEUDO-CODE FOR 1R

For each attribute,
 For each value of the attribute, make a rule:
 Count how often each class appears
 Find the most frequent class
 Make the rule assign that class to
 .. this attribute-value
 Calculate the error rate of the rules
 Choose the rules with the smallest error rate

■ Note: “missing” is treated as a separate attribute value

27

112

EVALUATING THE WEATHER ATTRIBUTES

Attribute	Rules	Errors	Total errors
Outlook	Sunny → No	2/5	4/14
	Overcast → Yes	0/4	
	Rainy → Yes	2/5	
Temp			
Humidity			
Windy			

28

112

EVALUATING THE WEATHER ATTRIBUTES (CONT.)

Attribute	Rules	Errors	Total errors
Outlook	Sunny → No	2/5	4/14
	Overcast → Yes	0/4	
	Rainy → Yes	2/5	
Temp	Hot → No*	2/4	5/14
	Mild → Yes	2/6	
	Cool → Yes	1/4	
Humidity	High → No	3/7	4/14
	Normal → Yes	1/7	
Windy	False → Yes	2/8	5/14
	True → No*	3/6	

* indicates a tie

29

112

OUTPUT OF ONER

Outlook Sunny → No
Overcast → Yes
Rainy → Yes

30

112

DEALING WITH NUMERIC ATTRIBUTES

- Discretizing Numeric Attributes
- Divide each attribute's range into intervals.
 - ▶ Sort instances according to the attribute's values.
 - ▶ Place breakpoints where the class changes (the majority class).
 - ▶ This minimizes the total error.

31

112

WEATHER DATA WITH SOME NUMERIC ATTRIBUTES

Outlook	Temperature	Humidity	Windy	Play
sunny	85	85	false	no
sunny	80	90	true	no
overcast	83	86	false	yes
rainy	70	96	false	yes
rainy	68	80	false	yes
rainy	65	70	true	no
overcast	64	65	true	yes
sunny	72	95	false	no
sunny	69	70	false	yes
rainy	75	80	false	yes
sunny	75	70	true	yes
overcast	72	90	true	yes
overcast	81	75	false	yes
rainy	71	91	true	no

32

112

Outlook	Temperature	Humidity	Windy	Play
sunny	85	85	false	no
sunny	80	90	true	no
overcast	83	86	false	yes
rainy	70	96	false	yes
rainy	68	80	false	yes
rainy	65	70	true	no
overcast	64	65	true	yes
sunny	72	95	false	no
sunny	69	70	false	yes
rainy	75	80	false	yes
sunny	75	70	true	yes
overcast	72	90	true	yes
overcast	81	75	false	yes
rainy	71	91	true	no

■ Example: temperature from weather data

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

33

112

THE PROBLEM OF OVERFITTING

- This procedure is very sensitive to noise.
 - One instance with an incorrect class label will probably produce a separate interval.
- Simple solution: enforce minimum number of instances in majority class per interval.

34

112

DISCRETIZATION EXAMPLE

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

Example (with min = 3):

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

Final result for temperature attribute

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

35

112

EVALUATING THE WEATHER ATTRIBUTES (CONT.)

Attribute	Rules	Errors	Total errors
Outlook	Sunny → No	2/5	4/14
	Overcast → Yes	0/4	
	Rainy → Yes	2/5	
Temp	≤ 77.5 → Yes	3/10	5/14
	> 77.5 → No*	2/4	
Humidity	≤ 82.5 → Yes	1/7	3/14
	≥ 82.5 and ≤ 95.5 → No	2/6	
	≤ 95.5 → Yes	0/1	
Windy	False → Yes	2/8	5/14
	True → No*	3/6	

¹* indicates a tie

36

112

Humidity $\leq 82.5 \rightarrow \text{Yes}$
 > 82.5 and $\leq 95.5 \rightarrow \text{No}$
 $> 95.5 \rightarrow \text{Yes}$

DECISION TREES

DECISION TREES

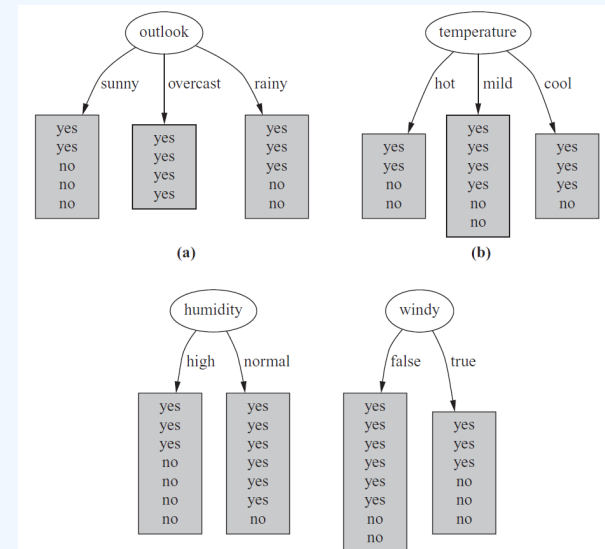
What is a Decision Tree?

Learns a tree-like structure, where each node represents an attribute (or a test) and each branch represents a possible value (or outcome of the test). The leaves of the tree represent the final class.

How to Build a Tree Automatically from Data

- Divide and conquer approach (Top-Down Induction of Decision Trees).
- Select an attribute to place as the root node.
- Make one branch for each possible value.
- Repeat the process for each branch.
- Stop if all instances at a node have the same classification.

WHICH ATTRIBUTE TO SELECT?



CRITERION FOR ATTRIBUTE SELECTION

- **Objective:** Find the best attribute to create the smallest decision tree.
- **Heuristic:** Choose the attribute that produces the "purest" nodes.
- **Popular Selection Criterion:** Information Gain.
 - ▶ Information gain increases with the average purity of the subsets.
- **Strategy:** Amongst attributes available for splitting, choose the attribute that gives the greatest information gain.
- **Information Gain Requirement:** A measure of impurity.
- **Impurity Measure:** Entropy of the class distribution (from information theory).

40

112

COMPUTING INFORMATION

- We have a probability distribution: the class distribution in a subset of instances.
- The expected information required to determine an outcome (i.e., class value) is the distribution's entropy.
- Formula for computing the entropy:

$$\text{Entropy}(p_1, p_2, \dots, p_n) = -p_1 \log p_1 - p_2 \log p_2 \cdots - p_n \log p_n$$

- Using base-2 logarithms, entropy gives the information required in expected bits.
- Entropy is maximal when all classes are equally likely and minimal when one of the classes has probability 1.

41

112

INFORMATION GAIN

- Expected Information

$$\text{Info}(D) = - \sum_{i=1}^m p_i \log_2 p_i$$

- Expected information after partitioned by attribute A

$$\text{Info}_A(D) = \sum_{i=1}^v \frac{|D_i|}{|D|} \times \text{Info}(D_i)$$

- Information Gain

$$\text{Gain}(A) = \text{Info}(D) - \text{Info}_A(D)$$

42

112

EXAMPLE: ATTRIBUTE OUTLOOK

- Outlook = Sunny:

$$\begin{aligned} \text{Info}([2, 3]) &= \text{entropy}\left(\frac{2}{5}, \frac{3}{5}\right) = -\frac{2}{5} \log_2 \left(\frac{2}{5}\right) - \frac{3}{5} \log_2 \left(\frac{3}{5}\right) \\ &= 0.971 \text{ bits} \end{aligned}$$

- Outlook = Overcast: $\text{Info}([4, 0]) = 0.0$ bits

$$\begin{aligned} \text{Info}([4, 0]) &= \text{entropy}(1, 0) = -1 \log_2(1) - 0 \log_2(0)^1 \\ &= 0.0 \text{ bits} \end{aligned}$$

- Outlook = Rainy: $\text{Info}([3, 2]) = 0.971$ bits

- Expected information for attribute:

$$\begin{aligned} \text{info}([2, 3], [4, 0], [3, 2]) &= \frac{5}{14} \times 0.971 + \frac{4}{14} \times 0 + \frac{5}{14} \times 0.971 \\ &= 0.693 \text{ bits} \end{aligned}$$

¹Note: $\log(0)$ is not defined, but we evaluate $0 \cdot \log(0)$ as zero

43

112

COMPUTING INFORMATION GAIN

- Information gain = information before splitting - information after splitting.

- Example:

$$\begin{aligned}\text{Gain}(\text{Outlook}) &= \text{Info}([9, 5]) - \text{info}([2, 3], [4, 0], [3, 2]) \\ &= 0.940 - 0.693 \\ &= 0.247 \text{ bits}\end{aligned}$$

- Information gain for attributes from weather data:

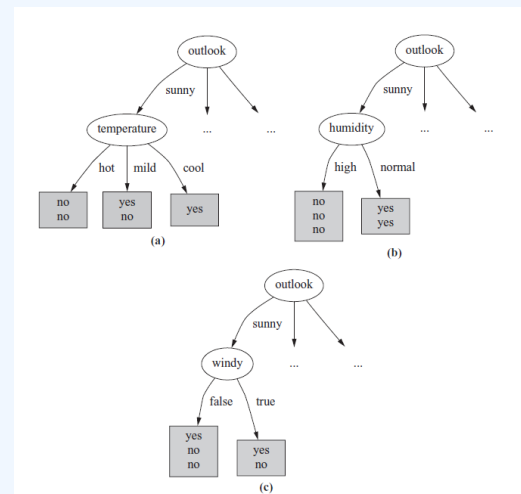
$$\begin{aligned}\text{Gain}(\text{Outlook}) &= 0.247 \text{ bits} \\ \text{Gain}(\text{Temperature}) &= 0.029 \text{ bits} \\ \text{Gain}(\text{Humidity}) &= 0.152 \text{ bits} \\ \text{Gain}(\text{Windy}) &= 0.048 \text{ bits}\end{aligned}$$

- Select the attribute with the highest gain ratio
- Problems: highly branching attributes in Information Gain, Example, ID code

44

112

CONTINUING TO SPLIT

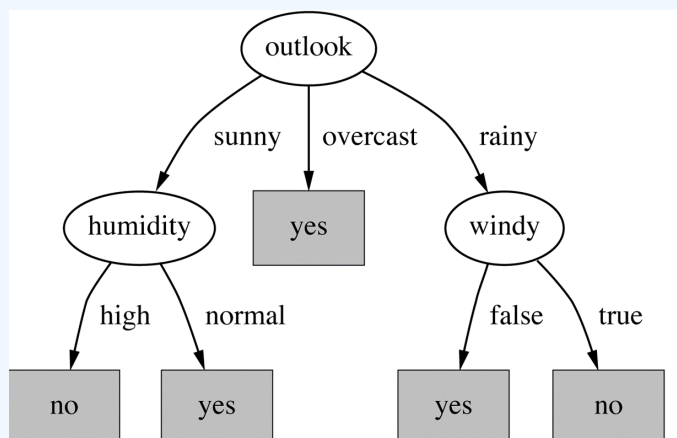


- Gain(Temperature) = 0.571 bits
- Gain(Humidity) = 0.971 bits
- Gain(Windy) = 0.020 bits

45

112

FINAL DECISION TREE




- Note: not all leaves need to be pure; sometimes identical instances have different classes
 - Splitting stops when data cannot be split any further

46

112

HIGHLY-BRANCHING ATTRIBUTES

- Problematic: attributes with a large number of values (extreme case: ID code)
- Subsets are more likely to be pure if there is a large number of values
 - Information gain is biased towards choosing attributes with a large number of values
 - This may result in overfitting (selection of an attribute that is non-optimal for prediction)



- In the above, all (single-instance) subsets have entropy zero!
- This means the information gain is maximal for this ID code attribute

47

112

GAIN RATIO

- Gain ratio is a modification of the information gain that reduces its bias towards attributes with many values.
- Gain ratio takes number and size of branches into account when choosing an attribute.
- It corrects the information gain by taking the intrinsic information of a split into account.
- Intrinsic information: entropy of the distribution of instances into branches.
- Measures how much info do we need to tell which branch a randomly chosen instance belongs to.

48

112

COMPUTING THE GAIN RATIO

- Example: intrinsic information of ID code

$$\frac{1}{14}(\text{info}([0, 1]) + \text{info}([0, 1]) + \text{info}([1, 0]) + \dots + \text{info}([1, 0]) + \text{info}([0, 1]))$$

- Value of attribute should decrease as intrinsic information gets larger.
- The gain ratio is defined as the information gain of the attribute divided by its intrinsic information.
- Example (outlook at root node):
 - ▶ Gain: $0.940 - 0.693 = 0.247$
 - ▶ Split info: $\text{info}([5, 4, 5]) = 1.577$
 - ▶ Gain ratio: $0.247 / 1.577 = 0.156$

49

112

ALL GAIN RATIOS FOR THE WEATHER DATA

Outlook	Temperature
Info: 0.693 Gain: $0.940 - 0.693 = 0.247$ Split info: $\text{info}([5, 4, 5]) = 1.577$ Gain ratio: $0.247 / 1.577 = 0.157$	Info: 0.911 Gain: $0.940 - 0.911 = 0.029$ Split info: $\text{info}([4, 6, 4]) = 1.557$ Gain ratio: $0.029 / 1.557 = 0.019$
Humidity	Windy
Info: 0.788 Gain: $0.940 - 0.788 = 0.152$ Split info: $\text{info}([7, 7]) = 1.000$ Gain ratio: $0.152 / 1 = 0.152$	Info: 0.892 Gain: $0.940 - 0.892 = 0.048$ Split info: $\text{info}([8, 6]) = 0.985$ Gain ratio: $0.048 / 0.985 = 0.049$

50

112

MORE ON THE GAIN RATIO

- “Outlook” still comes out top.
- However: “ID code” has greater gain ratio.
- Standard fix: ad hoc test to prevent splitting on that type of identifier attribute.
- Problem with gain ratio: it may overcompensate.
- May choose an attribute just because its intrinsic information is very low.
- Standard fix: only consider attributes with greater than average information gain.
- Both tricks are implemented in the well-known C4.5 decision tree learner.

51

112

- The **Gini** index is used in CART.

$$Gini(D) = 1 - \sum_{i=1}^m p_i^2 \quad (1)$$

- It considers a binary split for each attribute. **Gini** after partitioned by attribute A:

$$Gini_A(D) = \frac{|D_1|}{|D|} Gini(D_1) + \frac{|D_2|}{|D|} Gini(D_2) \quad (2)$$

- The reduction of impurity:

$$\Delta Gini(A) = Gini(D) - Gini_A(D) \quad (3)$$

- 1. Handling Numeric Attributes
- 2. Handling Missing Values
- 3. Generalizing the Model

- Extending ID3:
 - ▶ to permit numeric attributes: *straightforward*
 - ▶ to deal sensibly with missing values: *trickier*
 - ▶ stability for noisy data: *requires pruning mechanism*
 - **Problem:** ID3 performs well on the training set and does not generalize well on the independent test sets
- End result: C4.5 (Quinlan)
 - ▶ Best-known and (probably) most widely-used learning algorithm
 - ▶ Commercial successor: C5.0 (now available freely under GNU General Public License)

- **Standard method:** binary splits
 - ▶ E.g. temp < 45
- Unlike nominal attributes, every attribute has many possible split points
- **Solution is straightforward extension:**
 - ▶ Evaluate information gain (or other measure) for every possible split point of attribute
 - ▶ Choose “best” split point
 - ▶ Information gain for best split point is info gain for attribute
- Computationally more demanding

WEATHER DATA WITH SOME NUMERIC ATTRIBUTES

Outlook	Temperature	Humidity	Windy	Play
sunny	85	85	false	no
sunny	80	90	true	no
overcast	83	86	false	yes
rainy	70	96	false	yes
rainy	68	80	false	yes
rainy	65	70	true	no
overcast	64	65	true	yes
sunny	72	95	false	no
sunny	69	70	false	yes
rainy	75	80	false	yes
sunny	75	70	true	yes
overcast	72	90	true	yes
overcast	81	75	false	yes
rainy	71	91	true	no

56

112

EXAMPLE

- Split on temperature attribute:

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

- E.g. temperature < 71.5: yes/4, no/2
temperature ≥ 71.5: yes/5, no/3
- Expected information for temperature attribute:

$$\text{info}([4, 2], [5, 3]) = \frac{6}{14} \times \text{Info}([4, 2]) + \frac{8}{14} \times \text{Info}([5, 3]) \\ = 0.939 \text{ bits}$$

- Place split points halfway between values
- Can evaluate all split points in one pass!

57

112

MISSING VALUES

- Treat missing values as another possible value of the attribute (if significant)
- Ignore all the instances with missing values (lose information)
- Split instances with missing values into pieces:
 - A piece going down a branch receives a weight proportional to the popularity of the branch: weights sum to 1
 - Info gain works with fractional instances: use sums of weights instead of counts
 - During classification, split the instance into pieces in the same way: merge probability distribution using weights

58

112

PRUNING

- Prevent overfitting to noise in the data
- “Prune” the decision tree
- Two strategies:
 - Postpruning (backward pruning):** take a fully-grown decision tree and discard unreliable parts
 - Prepruning (forward pruning):** stop growing a branch when information becomes unreliable (decide on when to stop developing subtrees)
- Postpruning preferred in practice – prepruning can “stop early”

59

112

PREPRUNING

- Based on statistical significance test
 - ▶ Stop growing the tree when there is no *statistically significant association between any attribute and the class* at a particular node
- Pre-pruning may stop the growth process prematurely: *early stopping*
- Prepruning is faster than postpruning

60

112

POSTPRUNING

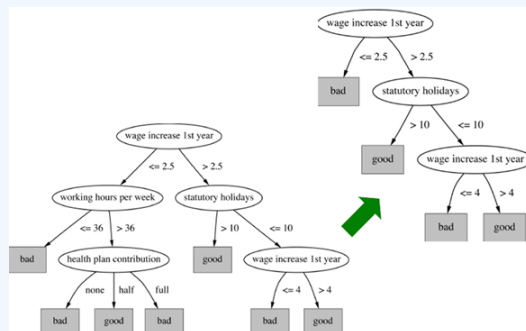
- First, build full tree and then, prune it
 - ▶ Fully-grown tree shows all attribute interactions
- Problem: some subtrees might be due to chance effects
- Two pruning operations:
 - ▶ Subtree replacement
 - ▶ Subtree raising
- Possible strategies:
 - ▶ error estimation
 - ▶ significance testing
 - ▶ MDL principle

61

112

SUBTREE REPLACEMENT


- The idea is to select some subtrees and replace them with single leaves.
- Bottom-up: Consider replacing a tree only after considering all its subtrees



62

112

SUBTREE RAISING



- Delete node
- Redistribute instances
- Slower than subtree replacement

63

112

DISCUSSION

- Top-down induction of decision trees: ID3, algorithm developed by Ross Quinlan.
- Gain ratio just one modification of this basic algorithm.
- C4.5 tree learner deals with numeric attributes, missing values, noisy data.
- Similar approach: CART tree learner.
- Uses Gini index rather than entropy to measure impurity.
- There are many other attribute selection criteria! (But little difference in accuracy of result).
 - ▶ Information Gain
 - ▶ Gain Ratio
 - ▶ Gini index (CART trees – used in scikit-learn)
 - ▶ Chi-Square

64

112

DECISION RULES (PRISM)

GENERATING RULES

- Decision tree can be converted into a rule set
- **Straightforward conversion:**
 - ▶ each path to the leaf becomes a rule – makes an overly complex rule set
- **More effective conversions are not trivial**
 - ▶ (e.g., C4.8 tests each node in root-leaf path to see if it can be eliminated without loss in accuracy)

65

112

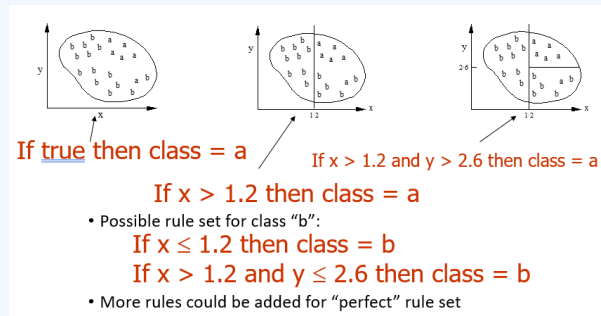
COVERING ALGORITHMS

- Strategy for generating a rule set directly: for each class in turn, find a rule set that covers all instances in it (excluding instances not in the class)
- This approach is called a **covering** approach because at each stage a rule is identified that covers some of the instances

66

112

EXAMPLE: GENERATING A RULE



67

112

A SIMPLE COVERING ALGORITHM


- Generates a rule by adding tests that maximize the rule's accuracy
- Similar to situation in decision trees: problem of selecting an attribute to split on
 - But: decision tree inducer maximizes overall purity
- Each new test reduces the rule's coverage

68

112

A SIMPLE COVERING ALGORITHM

- Generates a rule by adding tests that maximize the rule's accuracy
- Similar to situation in decision trees: problem of selecting an attribute to split on
 - But: decision tree inducer maximizes overall purity
- Each new test reduces the rule's coverage



69

112

SELECTING A TEST

- **Goal:** maximize accuracy
 - t = total number of instances covered by rule
 - p = positive examples of the class covered by rule
 - $t - p$ = number of errors made by rule
 - \Rightarrow Select test that maximizes the ratio p/t
- We are finished when $p/t = 1$ or the set of instances can't be split any further

70

112

CONTACT LENS DATA

Age	Spectacle Prescription	Astigmatism	Tear Production Rate	Recommended Lenses
young	myope	no	reduced	none
young	myope	no	normal	soft
young	myope	yes	reduced	none
young	myope	yes	normal	hard
young	hypermetrope	no	reduced	none
young	hypermetrope	no	normal	soft
young	hypermetrope	yes	reduced	none
young	hypermetrope	yes	normal	hard
pre-presbyopic	myope	no	reduced	none
pre-presbyopic	myope	no	normal	soft
pre-presbyopic	myope	yes	reduced	none
pre-presbyopic	myope	yes	normal	hard
pre-presbyopic	hypermetrope	no	reduced	none
pre-presbyopic	hypermetrope	no	normal	soft
pre-presbyopic	hypermetrope	yes	reduced	none
pre-presbyopic	hypermetrope	yes	normal	none
presbyopic	myope	no	reduced	none
presbyopic	myope	no	normal	none
presbyopic	myope	yes	reduced	none
presbyopic	myope	yes	normal	hard
presbyopic	hypermetrope	no	reduced	none
presbyopic	hypermetrope	no	normal	soft
presbyopic	hypermetrope	yes	reduced	none
presbyopic	hypermetrope	yes	normal	none

71

112

EXAMPLE: CONTACT LENS DATA, 1

■ Rule we seek:

If ?
then recommendation = hard

■ Possible tests:

- ▶ Age = Young
- ▶ Age = Pre-presbyopic
- ▶ Age = Presbyopic
- ▶ Spectacle prescription = Myope
- ▶ Spectacle prescription = Hypermetrope
- ▶ Astigmatism = no
- ▶ Astigmatism = yes
- ▶ Tear production rate = Reduced
- ▶ Tear production rate = Normal

2/8

72

112

EXAMPLE: CONTACT LENS DATA, 2

■ Rule we seek:

If ?
then recommendation = hard

■ Possible tests:

- ▶ Age = Young 2/8
- ▶ Age = Pre-presbyopic 1/8
- ▶ Age = Presbyopic 1/8
- ▶ Spectacle prescription = Myope 3/12
- ▶ Spectacle prescription = Hypermetrope 1/12
- ▶ Astigmatism = no 0/12
- ▶ Astigmatism = yes 4/12
- ▶ Tear production rate = Reduced 0/12
- ▶ Tear production rate = Normal 4/12

73

112

MODIFIED RULE AND RESULTING DATA

■ Rule with best test added:

If astigmatism = yes
then recommendation = hard

■ Instances covered by modified rule:

Age	Spectacle prescription	Astigmatism	Tear production rate	Recommended lenses
Young	Myope	Yes	Reduced	None
Young	Myope	Yes	Normal	Hard
Young	Hypermetrope	Yes	Reduced	None
Young	Hypermetrope	Yes	Normal	Hard
Pre-presbyopic	Myope	Yes	Reduced	None
Pre-presbyopic	Myope	Yes	Normal	Hard
Pre-presbyopic	Hypermetrope	Yes	Reduced	None
Pre-presbyopic	Hypermetrope	Yes	Normal	None
Presbyopic	Myope	Yes	Reduced	None
Presbyopic	Myope	Yes	Normal	Hard
Presbyopic	Hypermetrope	Yes	Reduced	None
Presbyopic	Hypermetrope	Yes	Normal	None

74

112

FURTHER REFINEMENT, 1

■ Current state:

```
If astigmatism = yes
and ?
then recommendation = hard
```

■ Possible tests:

- ▶ Age = Young
- ▶ Age = Pre-presbyopic
- ▶ Age = Presbyopic
- ▶ Spectacle prescription = Myope
- ▶ Spectacle prescription = Hypermetrope
- ▶ Tear production rate = Reduced
- ▶ Tear production rate = Normal

2/4

75

112

FURTHER REFINEMENT, 2

■ Current state:

```
If astigmatism = yes
and ?
then recommendation = hard
```

■ Possible tests:

- ▶ Age = Young
- ▶ Age = Pre-presbyopic
- ▶ Age = Presbyopic
- ▶ Spectacle prescription = Myope
- ▶ Spectacle prescription = Hypermetrope
- ▶ Tear production rate = Reduced
- ▶ Tear production rate = Normal

2/4

1/4

1/4

3/6

1/6

0/6

4/6

76

112

MODIFIED RULE AND RESULTING DATA

■ Rule with best test added:

```
If astigmatism = yes
and tear production rate = normal
then recommendation = hard
```

■ Instances covered by modified rule:

Age	Spectacle prescription	Astigmatism	Tear production rate	Recommended lenses
Young	Myope	Yes	Normal	Hard
Young	Hypermetrope	Yes	Normal	hard
Pre-presbyopic	Myope	Yes	Normal	Hard
Pre-presbyopic	Hypermetrope	Yes	Normal	None
Presbyopic	Myope	Yes	Normal	Hard
Presbyopic	Hypermetrope	Yes	Normal	None

77

112

FURTHER REFINEMENT, 3

■ Current state:

```
If astigmatism = yes
and tear production rate = normal
and ?
then recommendation = hard
```

■ Possible tests:

- ▶ Age = Young
- ▶ Age = Pre-presbyopic
- ▶ Age = Presbyopic
- ▶ Spectacle prescription = Myope
- ▶ Spectacle prescription = Hypermetrope

78

112

FURTHER REFINEMENT, 4

■ Current state:

```
If astigmatism = yes
  and tear production rate = normal
  and ?
  then recommendation = hard
```

■ Possible tests:

- ▶ Age = Young 2/2
- ▶ Age = Pre-presbyopic 1/2
- ▶ Age = Presbyopic 1/2
- ▶ Spectacle prescription = Myope 3/3
- ▶ Spectacle prescription = Hypermetrope 1/3

■ Tie between the first and the fourth test

- ▶ We choose the one with greater coverage

79

112

THE RESULT

■ Final rule:

```
If astigmatism = yes
  and tear production rate = normal
  and spectacle prescription = myope
  then recommendation = hard
```

■ Second rule for recommending “hard lenses”: (built from instances not covered by the first rule)

```
If age = young and astigmatism = yes
  and tear production rate = normal
  then recommendation = hard
```

■ These two rules cover all “hard lenses”:

- ▶ Process is repeated with other two classes

80

112

PSEUDO-CODE FOR PRISM

```
For each class C
  Initialize E to the instance set
  While E contains instances in class C
    Create a rule R with an empty left-hand side that predicts class C
    Until R is perfect (or there are no more attributes to use) do
      For each attribute A not mentioned in R, and each value v,
        Consider adding the condition A = v to the left-hand side of R
        Select A and v to maximize the accuracy p/t
        (break ties by choosing the condition with the largest p)
      Add A = v to R
    Remove the instances covered by R from E
```

81

112

SEPARATE AND CONQUER ALGORITHMS

■ PRISM is a separate-and-conquer algorithm

- ▶ Identify a rule that covers many instances in the class
- ▶ Separate out the covered instances
- ▶ Continue with the remaining instances

82

112

NEAREST NEIGHBOR

INSTANCE-BASED REPRESENTATION

- Simplest form of learning: **rote learning**
 - ▶ Training instances are searched for instance that most closely resembles new instance
 - ▶ The instances themselves represent the knowledge
 - ▶ Also called *instance-based learning*
- **Similarity function** defines what's "learned"
- Instance-based learning is *lazy learning*
- Methods:
 - ▶ **nearest-neighbor**
 - ▶ **k-nearest-neighbor**

83

112

THE DISTANCE FUNCTION

- Simplest case: **one numeric attribute**
 - ▶ Distance is the difference between the two attribute values involved (or a function thereof)
- **Several numeric attributes**: normally, Euclidean distance is used and attributes are normalized
- **Nominal attributes**: distance is set to 1 if values are different, 0 if they are equal
- **Are all attributes equally important?**
 - ▶ Weighting the attributes might be necessary

84

112

INSTANCE-BASED LEARNING: NEAREST NEIGHBOR ALGORITHM

- Distance function defines what's learned
- Most instance-based schemes use **Euclidean distance**:

$$\sqrt{(a_1^{(1)} - a_1^{(2)})^2 + (a_2^{(1)} - a_2^{(2)})^2 + \dots + (a_k^{(1)} - a_k^{(2)})^2}$$

- $a^{(1)}$ and $a^{(2)}$: two instances with k attributes
- Taking the square root is not required when comparing distances
- Other popular metric: **city-block (Manhattan) metric**
 - ▶ Adds differences without squaring them:

$$|a_1^{(1)} - a_1^{(2)}| + |a_2^{(1)} - a_2^{(2)}| + \dots + |a_k^{(1)} - a_k^{(2)}|$$

85

112

DISTANCE FUNCTION (CONT.)

■ Minkowski distance

$$\left(\sum_{i=1}^k |a_i^{(1)} - a_i^{(2)}|^p \right)^{\frac{1}{p}}$$

- When $p = 1$, Manhattan distance
- When $p = 2$, Euclidean distance

86

112

EXAMPLE: DISSIMILARITY MATRICES

■ Points and Attributes:

Point	Attribute 1	Attribute 2
x1	1	2
x2	3	5
x3	2	0
x4	4	5

Dissimilarity Matrices: Manhattan

L	x1	x2	x3	x4
x1	0	5	3	6
x2	5	0	6	1
x3	3	6	0	7
x4	6	1	7	0

(L₁):

Euclidean

(L₂):


L2	x1	x2	x3	x4
x1	0	3.61	2.24	4.24
x2	3.61	0	5.10	1
x3	2.24	5.10	0	5.39
x4	4.24	1	5.39	0

87

112

EXAMPLE

point	attribute 1	attribute 2
x1	1	2
x2	3	5
x3	2	0
x4	4	5



Dissimilarity Matrices

Manhattan (L₁)

L	x1	x2	x3	x4
x1	0			
x2		0		
x3			0	
x4				0

Euclidean (L₂)


L2	x1	x2	x3	x4
x1	0			
x2		0		
x3			0	
x4				0

88

112

EXAMPLE

point	attribute 1	attribute 2
x1	1	2
x2	3	5
x3	2	0
x4	4	5



Dissimilarity Matrices

Manhattan (L₁)

L	x1	x2	x3	x4
x1	0			
x2	5	0		
x3	3	6	0	
x4	6	1	7	0


Euclidean (L₂)

L2	x1	x2	x3	x4
x1	0			
x2	3.61	0		
x3	2.24	5.1	0	
x4	4.24	1	5.39	0

89

112

NEAREST NEIGHBOR CLASSIFICATION



90

112

NORMALIZATION AND OTHER ISSUES

- Different attributes are measured on different scales → need to be **normalized**:

$$a_i = \frac{v_i - \min v_i}{\max v_i - \min v_i} \quad \text{or} \quad a_i = \frac{v_i - \text{Avg}(v_i)}{\text{StDev}(v_i)}$$

- v_i : the actual value of attribute i
- Nominal attributes: distance either 0 or 1
- Common policy for **missing values**: assumed to be maximally distant (given normalized attributes)

91

112

DISCUSSION OF NEAREST NEIGHBOR


- Often very accurate
- ... but slow:
 - ▶ Simple version scans entire training data to derive a prediction
 - ▶ Remedy: use of kD-tree and Ball-trees
- Assumes all attributes are equally important
 - ▶ Remedy: attribute selection or weights
- Possible remedies against noisy instances:
 - ▶ Take a majority vote over the k nearest neighbors
 - ▶ Removing noisy instances from dataset (difficult!)
- Statisticians have used k-NN since early 1950s

92

112

STATISTICAL MODELING

- “Opposite” of 1R: use all the attributes
- Two assumptions: Attributes are
 - ▶ *equally important*
 - ▶ *statistically independent* (given the class value)
 - I.e., knowing the value of one attribute says nothing about the value of another (if the class is known)
- Independence assumption is almost never correct!
- But ... this scheme works well in practice



- Naïve Bayes works surprisingly well (even if independence assumption is clearly violated)
- Why? Because classification doesn't require accurate probability estimates *as long as* maximum probability is assigned to the correct class
- However: adding too many redundant attributes will cause problems (e.g., identical attributes)
- Note also: many numeric attributes are not normally distributed (→ *kernel density estimators*)

Outlook			Temperature			Humidity		Windy		Play	
Yes		No	Yes		No	Yes		No	Yes		No
Sunny	2	3	Hot	2	2	High		False		9	5
Overcast	4	0	Mild	4	2	Normal		True			
Rainy	3	2	Cool	3	1						
Sunny	2/9	3/5	Hot			High		False		9/14	5/14
Overcast	4/9	0/5	Mild			Normal		True			
Rainy	3/9	2/5	Cool								

PROBABILITIES FOR WEATHER DATA

Outlook			Temperature			Humidity			Windy			Play	
			Yes	No		Yes	No		Yes	No		Yes	No
Sunny	2	3	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	0	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	2	Cool	3	1								
Sunny	2/9	3/5	Hot	2/9	2/5	High	3/9	4/5	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	Mild	4/9	2/5	Normal	6/9	1/5	True	3/9	3/5		
Rainy	3/9	2/5	Cool	3/9	1/5								

97

112

PROBABILITIES FOR WEATHER DATA

■ A new day:

Outlook	Temp.	Humidity	Windy	Play
Sunny	Cool	High	True	?

Likelihood of the two classes:

■ For “yes” = $\frac{2}{9} \times \frac{3}{9} \times \frac{3}{9} \times \frac{3}{9} \times \frac{9}{14} = 0.0053$

■ For “no” = $\frac{3}{5} \times \frac{1}{5} \times \frac{4}{5} \times \frac{3}{5} \times \frac{5}{14} = 0.0206$

Conversion into a probability by normalization:

■ $P(\text{“yes”}) = \frac{0.0053}{0.0053+0.0206} = 0.205$

■ $P(\text{“no”}) = \frac{0.0206}{0.0053+0.0206} = 0.795$

98

112

BAYES’S RULE

■ Probability of event H given evidence E :

$$\Pr(H | E) = \frac{\Pr(E | H) \cdot \Pr(H)}{\Pr(E)}$$

■ **A priori** probability of H :

► Probability of event *before* evidence is seen ($\Pr(H)$)

■ **A posteriori** probability of H :

► Probability of event *after* evidence is seen ($\Pr(H | E)$)

99

112

WEATHER DATA EXAMPLE

Evidence E :

Outlook	Temp.	Humidity	Windy	Play
Sunny	Cool	High	True	?

Probability of class “yes”:

$$\begin{aligned}
 \Pr(\text{yes} | E) &= \Pr[\text{Outlook} = \text{Sunny} | \text{yes}] \\
 &\quad \times \Pr[\text{Temperature} = \text{Cool} | \text{yes}] \\
 &\quad \times \Pr[\text{Humidity} = \text{High} | \text{yes}] \\
 &\quad \times \Pr[\text{Windy} = \text{True} | \text{yes}] \\
 &\quad \times \frac{\Pr[\text{yes}]}{\Pr[E]} \\
 &= \frac{\frac{2}{9} \times \frac{3}{9} \times \frac{3}{9} \times \frac{3}{9} \times \frac{9}{14}}{\Pr[E]}
 \end{aligned}$$

100

112

PREDICT THE CLASS VALUE OF THE FOLLOWING DAY

A new day:

Outlook	Temp.	Humidity	Windy	Play
Overcast	Cool	High	True	?

101

112

THE “ZERO-FREQUENCY PROBLEM”

- What if an attribute value doesn't occur with every class value?
 - ▶ e.g., “Outlook = Overcast” for class “no”
 - ▶ Probability will be zero! $\Pr(\text{Outlook} = \text{Overcast} \mid \text{yes}) = 0$
 - ▶ A *posteriori* probability will also be zero! $\Pr(\text{yes} \mid E) = 0$
 - ▶ (No matter how likely the other values are!)
- Remedy: add 1 to the count for every attribute value–class combination (**Laplace estimator**)
- Result: probabilities will never be zero! (Also: stabilizes probability estimates)

102

112

LAPLACE ESTIMATOR FOR WEATHER DATA

Outlook	Temperature		Humidity		Windy		Play	
	Yes	No	Yes	No	Yes	No	Yes	No
Sunny	2	3 + 1	Hot	2	2	High	3	4
Overcast	4	0 + 1	Mild	4	2	Normal	6	1
Rainy	3	2 + 1	Cool	3	1	False	6	2
						True	3	3
Sunny	2/9	4/8	Hot	2/9	2/5	High	3/9	4/5
Overcast	4/9	1/8	Mild	4/9	2/5	Normal	6/9	1/5
Rainy	3/9	3/8	Cool	3/9	1/5	False	6/9	2/5
						True	3/9	3/5
							9/14	5/14

103

112

MISSING VALUES

- **Training:** instance is not included in frequency count for attribute value–class combination
- **Classification:** attribute will be omitted from calculation
- **Example:**

Outlook	Temp.	Humidity	Windy	Play
?	Cool	High	True	?

$$\text{Likelihood of class "yes":} = \frac{3}{9} \times \frac{3}{9} \times \frac{3}{9} \times \frac{9}{14} = 0.0238$$

$$\text{Likelihood of class "no":} = \frac{1}{5} \times \frac{4}{5} \times \frac{3}{5} \times \frac{5}{14} = 0.0343$$

$$P(\text{yes}) = \frac{0.0238}{0.0238 + 0.0343} = 0.41 \quad (41\%)$$

$$P(\text{no}) = \frac{0.0343}{0.0238 + 0.0343} = 0.59 \quad (59\%)$$

104

112

CLASSIFYING A NEW DAY

■ A new day:

Outlook	Temp.	Humidity	Windy	Play
Sunny	66	90	true	?

105

112

NUMERIC ATTRIBUTES

- Usual assumption: attributes have a **normal** or **Gaussian** probability distribution (given the class)
- The *probability density function* for the normal distribution is defined by two parameters:
 - ▶ Sample mean μ :

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

- ▶ Standard deviation σ :

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)^2}$$

- Then the density function $f(x)$ is:

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

106

112

STATISTICS FOR WEATHER DATA

Outlook			Temperature		Humidity		Windy			Play	
			Yes	No	Yes	No	Yes	No		Yes	No
Sunny	2	3	64, 68,	65, 71,	65, 70,	70, 85,	False	6	2	9	5
Overcast	4	0	69, 70,	72, 80,	70, 75,	90, 91,	True	3	3		
Rainy	3	2	72, ...	85, ...	80, ...	95, ...					
Sunny	2/9	3/5	$\mu = 73$	$\mu = 75$	$\mu = 79$	$\mu = 86$	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	$\sigma = 6.2$	$\sigma = 7.9$	$\sigma = 10.2$	$\sigma = 9.7$	True	3/9	3/5		

Example density value:

$$f(\text{temperature} = 66 \mid \text{yes}) = \frac{1}{\sqrt{2\pi}6.2} e^{-\frac{(66-73)^2}{2 \times 6.2^2}} = 0.0340$$

$$f(\text{humidity} = 90 \mid \text{yes}) = \frac{1}{\sqrt{2\pi}10.2} e^{-\frac{(90-79)^2}{2 \times 10.2^2}} = 0.0221$$

107

112

CLASSIFYING A NEW DAY

■ A new day:

Outlook	Temp.	Humidity	Windy	Play
Sunny	66	90	true	?

$$\text{Likelihood of "yes"} = \frac{2}{9} \times 0.0340 \times 0.0221 \times \frac{3}{9} \times \frac{9}{14} = 0.000036$$

$$\text{Likelihood of "no"} = \frac{3}{5} \times 0.0291 \times 0.0380 \times \frac{3}{5} \times \frac{5}{14} = 0.000136$$

$$P(\text{"yes"}) = \frac{0.000036}{0.000036 + 0.000136} = 20.9\%$$

$$P(\text{"no"}) = \frac{0.000136}{0.000036 + 0.000136} = 79.1\%$$

Note: Missing values during training are not included in calculation of mean and standard deviation.

108

112

PROBABILITY DENSITIES

- Probability densities $f(x)$ can be greater than 1; hence, they are not probabilities.
 - ▶ However, they must integrate to 1: the area under the probability density curve must be 1.
- Approximate relationship between probability and probability density can be stated as:

$$P(x - \varepsilon/2 \leq X \leq x + \varepsilon/2) \approx \varepsilon f(x)$$

assuming ε is sufficiently small.

- When computing likelihoods, we can treat densities just like probabilities.

109

112

MULTINOMIAL NAÏVE BAYES I

- Version of naïve Bayes used for document classification using *bag of words* model.
- n_1, n_2, \dots, n_k : number of times word i occurs in the document.
- P_1, P_2, \dots, P_k : probability of obtaining word i when sampling from documents in class H .
- Probability of observing a particular document E given class H (based on *multinomial distribution*):

$$P(E | H) = N! \times \prod_{i=1}^k \frac{P_i^{n_i}}{n_i!}$$

- Note: this expression ignores the probability of generating a document of the right length.
 - ▶ This probability is assumed to be constant for all classes.

110

112

MULTINOMIAL NAÏVE BAYES II

- Suppose dictionary has two words, *yellow* and *blue*
- Suppose $P(\text{yellow} | H) = 75\%$ and $P(\text{blue} | H) = 25\%$
- Suppose E is the document “*blue yellow blue*”
- Probability of observing document:

$$P(\{\text{blue yellow blue}\} | H) = 3! \times \frac{0.75^1}{1!} \times \frac{0.25^2}{2!} = \frac{27}{64}$$

- Suppose there is another class H' with:

$$P(\text{yellow} | H') = 10\%, \quad P(\text{blue} | H') = 90\%$$

$$P(\{\text{blue yellow blue}\} | H') = 3! \times \frac{0.1^1}{1!} \times \frac{0.9^2}{2!} = \frac{243}{1000}$$

- Need to take prior probability of class into account using Bayes' rule
- Factorials do not need to be computed: they cancel out
- Underflows can be prevented by using logarithms

111

112

NAÏVE BAYES: DISCUSSION

- Naïve Bayes works surprisingly well (even if independence assumption is clearly violated)
- Why? Because classification doesn't require accurate probability estimates *as long as maximum probability is assigned to correct class*

112 / 112